

Foundations of Machine Learning

LINEAR REGRESSION (CONTD)

August 25, 2020

Recap: Least Squares Linear Regression

Given: Training data containing n samples where each $\mathbf{x}_i \in \mathbb{R}^d$, $\mathcal{D} = \{\mathbf{x}_i, y_i\}_{i=1}^n$

Least squares solution to linear regression is as follows:

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \underline{E(\mathbf{w}, \mathcal{D})} = \arg \min_{\mathbf{w}} \left\{ \sum_{i=1}^n \left(\sum_{j=0}^m w_j \phi_j(\mathbf{x}_i) - y_i \right)^2 \right\}$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \left(\underline{\|\Phi \mathbf{w} - \mathbf{y}\|^2} \right) = \underline{(\Phi^T \Phi)^{-1} \Phi^T \mathbf{y}}$$

where

$$\Phi = \begin{bmatrix} 1 & \phi_1(\mathbf{x}_1) & \dots & \phi_m(\mathbf{x}_1) \\ \dots & \dots & \dots & \dots \\ 1 & \phi_1(\mathbf{x}_n) & \dots & \phi_m(\mathbf{x}_n) \end{bmatrix}, \mathbf{w} = \begin{bmatrix} w_0 \\ \dots \\ w_m \end{bmatrix}, \mathbf{y} = \begin{bmatrix} y_1 \\ \dots \\ y_n \end{bmatrix}$$

$n \times (m+1)$ $(m+1) \times 1$ $n \times 1$

Probabilistic Modeling of Linear Regression

- From the training set $\mathcal{D} = \{\mathbf{x}_j, y_j\}_{j=1}^n$, assume each target y_j is drawn from a probabilistic model.

Probabilistic Modeling of Linear Regression

- From the training set $\mathcal{D} = \{\mathbf{x}_j, y_j\}_{j=1}^n$, assume each target y_j is drawn from a probabilistic model.
- Let us assume a noisy linear model parameterized by $\mathbf{w} \in \mathbb{R}^d$

$$y_j = \mathbf{w}^T \mathbf{x}_j + \epsilon_j$$

Each target is associated with an additive stochastic noise modeled as a random variable ϵ_j .

Probabilistic Modeling of Linear Regression

- From the training set $\mathcal{D} = \{\mathbf{x}_j, y_j\}_{j=1}^n$, assume each target y_j is drawn from a probabilistic model.
- Let us assume a noisy linear model parameterized by $\mathbf{w} \in \mathbb{R}^d$

$$y_j = \mathbf{w}^T \mathbf{x}_j + \epsilon_j$$

Each target is associated with an additive stochastic noise modeled as a random variable ϵ_j .

- Let us model the noise to come from a **Gaussian distribution**:

$$\epsilon_j \sim \mathcal{N}(\underbrace{0}_{\text{mean}}, \underbrace{\sigma^2}_{\text{variance}}).$$

Probabilistic Modeling of Linear Regression

- From the training set $\mathcal{D} = \{\mathbf{x}_j, y_j\}_{j=1}^n$, assume each target y_j is drawn from a probabilistic model.
- Let us assume a noisy linear model parameterized by $\mathbf{w} \in \mathbb{R}^d$

$$y_j = \mathbf{w}^T \mathbf{x}_j + \epsilon_j$$

Each target is associated with an additive stochastic noise modeled as a random variable ϵ_j .

- Let us model the noise to come from a **Gaussian distribution**:
 $\epsilon_j \sim \mathcal{N}(0, \sigma^2)$.
- Noise variables ϵ_j all have the same mean (0), same variance (σ^2) and $\text{Cov}(\epsilon_i, \epsilon_j) = 0$ for $i \neq j$.

Probabilistic Modeling of Linear Regression

- If $y_j = \mathbf{w}^T \mathbf{x}_j + \epsilon_j$ and $\epsilon_j \sim \mathcal{N}(0, \sigma^2)$, then each y_j is drawn from the following Gaussian:

$$y_j \sim \mathcal{N}(\mathbf{w}^T \mathbf{x}_j, \sigma^2)$$

Probabilistic Modeling of Linear Regression

- If $y_j = \mathbf{w}^T \mathbf{x}_j + \epsilon_j$ and $\epsilon_j \sim \mathcal{N}(0, \sigma^2)$, then each y_j is drawn from the following Gaussian:

$$y_j \sim \mathcal{N}(\mathbf{w}^T \mathbf{x}_j, \sigma^2)$$

- Objective: Estimate \mathbf{w} from this probabilistic model. We will look at **Maximum Likelihood Estimation (MLE)** in this lecture.

MLE, more generally

A set of independent and identically distributed (i.i.d.) observations

$\mathbf{y} = \{y_1, \dots, y_n\}$ are generated according to a probability model parameterized by

θ , i.e. $y_j \sim \underline{\text{Pr}(\mathbf{y}|\theta)}$

MLE, more generally

A set of independent and identically distributed (i.i.d.) observations $\mathbf{y} = \{y_1, \dots, y_n\}$ are generated according to a probability model parameterized by θ , i.e. $y_j \sim \Pr(\mathbf{y}|\theta)$

What is $\Pr(\mathbf{y}|\theta)$ (also referred to as the likelihood)?

MLE, more generally

A set of independent and identically distributed (i.i.d.) observations $\mathbf{y} = \{y_1, \dots, y_n\}$ are generated according to a probability model parameterized by θ , i.e. $y_j \sim \Pr(\mathbf{y}|\theta)$

What is $\Pr(\mathbf{y}|\theta)$ (also referred to as the *likelihood*)?

Find the parameters θ . What's one good way? Find θ s.t. it maximizes the likelihood of the observed data. Hence, **MLE**.

MLE contd

- Typically maximize log-likelihood i.e. $\log(\text{Pr}(\mathbf{y}|\theta))$ instead of the likelihood.
Would this affect the estimated parameters?

MLE contd

- Typically maximize log-likelihood i.e. $\log(\Pr(\mathbf{y}|\theta))$ instead of the likelihood.
Would this affect the estimated parameters?
- Why apply this log transformation?
 - Mathematical convenience
 - Numerical convenience

MLE contd

- Typically maximize log-likelihood i.e. $\log(\Pr(\mathbf{y}|\theta))$ instead of the likelihood.
Would this affect the estimated parameters?

- Why apply this log transformation?
- What is the log-likelihood?

$$\Pr(y_j; \theta) \text{ or } \mathcal{P}_{\theta}(y_j)$$

$$\mathcal{L}(\theta) = \log(\Pr(\mathbf{y}|\theta)) = \log\left(\prod_{j=1}^n \Pr(y_j|\theta)\right) = \sum_{j=1}^n \log(\Pr(y_j|\theta))$$

- The MLE estimate is given by:

$$\theta_{\text{MLE}} = \underset{\theta}{\operatorname{argmax}} \mathcal{L}(\theta) = \underset{\theta}{\operatorname{argmax}} \sum_{j=1}^n \log(\Pr(y_j|\theta))$$

Illustrating MLE using a simple coin toss example

Say we toss a coin N times. Each coin toss y_j is a binary random variable with a Bernoulli distribution $\Pr(y_j|\theta) = \underline{\theta}^{y_j}(1 - \theta)^{1-y_j}$. Estimate θ using MLE.

$$\begin{aligned} L(\theta) &= \log \prod_{j=1}^N P(y_j | \theta) = \sum_j \log P(y_j | \theta) \\ &= \sum_j y_j \log \theta + (1 - y_j) \log(1 - \theta) \end{aligned}$$

$$\frac{\partial L(\theta)}{\partial \theta} = 0 \Rightarrow \sum_j \frac{y_j}{\theta} - \frac{(1 - y_j)}{1 - \theta} = 0 \Rightarrow \boxed{\theta = \frac{\sum_j y_j}{N}}$$

Coming back to MLE for linear regression

Assuming that each target y_j is Gaussian-distributed, i.e.

$$\Pr(y_j|x_j, \mathbf{w}) = \mathcal{N}(\mathbf{w}^T x_j, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ \frac{-(y_j - \mathbf{w}^T x_j)^2}{2\sigma^2} \right\}$$

Likelihood of the data $\mathcal{L}(\mathbf{w}) = \Pr(\mathbf{y}|\mathbf{X}, \mathbf{w})$

Coming back to MLE for linear regression

Assuming that each target y_j is Gaussian-distributed, i.e.

$$\Pr(y_j|x_j, \mathbf{w}) = \mathcal{N}(\mathbf{w}^T x_j, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ \frac{-(y_j - \mathbf{w}^T x_j)^2}{2\sigma^2} \right\}$$

Likelihood of the data $\mathcal{L}(\mathbf{w}) = \Pr(\mathbf{y}|\mathbf{X}, \mathbf{w})$

$$\mathcal{L}(\mathbf{w}) = \log \prod_{j=1}^n \mathcal{P}(y_j | x_j, \mathbf{w}) = \sum_j \log \mathcal{P}(y_j | x_j, \mathbf{w})$$

MLE estimate: $\mathbf{w}_{\text{MLE}} = \underset{\mathbf{w}}{\operatorname{argmax}} \mathcal{L}(\mathbf{w})$

MLE estimate for probabilistic linear regression

To estimate \mathbf{w}_{MLE} , consider log-likelihood ($\log(\mathcal{L}(\mathbf{w}))$) instead of likelihood ($\mathcal{L}(\mathbf{w})$):

$$\log(\mathcal{L}(\mathbf{w})) = LL(\mathbf{w})$$

MLE estimate for probabilistic linear regression

To estimate \mathbf{w}_{MLE} , consider log-likelihood ($\log(\mathcal{L}(\mathbf{w}))$) instead of likelihood ($\mathcal{L}(\mathbf{w})$):

$$\log(\mathcal{L}(\mathbf{w})) = LL(\mathbf{w})$$

$$\begin{aligned} LL(\mathbf{w}) &= \sum_j \log P(y_j | x_j, \mathbf{w}) \\ &= \text{constant} - \sum_j \frac{(y_j - \mathbf{w}^T x_j)^2}{2\sigma^2} \end{aligned}$$

$$\mathbf{w}_{\text{MLE}} = \underset{\mathbf{w}}{\operatorname{argmax}} LL(\mathbf{w}) = \underset{\mathbf{w}}{\operatorname{argmin}} \sum_{j=1}^n (y_j - \mathbf{w}^T x_j)^2 \quad (\text{Same as least squares!})$$

Estimating \mathbf{w}

Instead of finding the least squares/MLE solution analytically, use a search algorithm that progressively decreases the error function $E(\mathbf{w}, \mathcal{D})$.

- Initialize \mathbf{w}
- Until we converge:

Find a new value of \mathbf{w} that reduces $E(\mathbf{w}, \mathcal{D})$

$$\mathbf{w}_{t+1} \leftarrow \mathbf{w}_t + \boxed{\phantom{\mathbf{w}_t}}$$

Estimating \mathbf{w}

Instead of finding the least squares/MLE solution analytically, use a search algorithm that progressively decreases the error function $E(\mathbf{w}, \mathcal{D})$.

- Initialize \mathbf{w}
- Until we converge:

Find a new value of \mathbf{w} that reduces $E(\mathbf{w}, \mathcal{D})$

Gradient descent is an iterative algorithm used to find the minimum of a function.

Gradient Descent

$$\mathbf{w} = \begin{bmatrix} w_1 \\ \vdots \\ w_t \\ \vdots \\ w_d \end{bmatrix}$$

Initialize $\mathbf{w} \leftarrow \vec{0}$ vector

Repeat until convergence:

$w_i \leftarrow w_i - \eta \frac{\partial E(\mathbf{w})}{\partial w_i}$ (Simultaneously update each w_i)

$$\|\nabla_{\mathbf{w}} E(\mathbf{w})\|_2 \leq \epsilon$$

$$\|E(\mathbf{w}_{t+1}) - E(\mathbf{w}_t)\|_2 \leq \epsilon$$

Compute $\frac{\partial E(\mathbf{w})}{\partial w_i}$

$$\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} E(\mathbf{w})$$

LEARNING
RATE

Using Gradient Descent (GD)

Weight update rule in GD: $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} E(\mathbf{w})$

$\sum_i E_i(\mathbf{w})$

Faster version: **Stochastic Gradient Descent (SGD)**

The weight update rule for SGD: $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} E(\mathbf{w}; \mathbf{x}_i, y_i)$ where the loss is computed for a single example randomly sampled from the training set.

Middle ground: **Mini-batch Gradient Descent**

The weight update rule becomes: $\mathbf{w} \leftarrow \mathbf{w} - \eta \nabla_{\mathbf{w}} E_B(\mathbf{w}; \{\mathbf{x}_i, y_i\}_{i=1}^{|B|})$ where the loss is computed for a batch of $|B|$ examples sampled from the training set.